



## Example: Disease Diagnosis



## General Approach for Building Classification Model



## k Nearest-Neighbor Classification

### Rule Learning

Day	Temperature	Outlook	Humidity	Windy	Play Golf?
07-05	hot	sunny	high	false	no
07-06	hot	sunny	high	true	no
07-07	hot	overcast	high	false	yes
07-09	cool	rain	normal	false	yes
07-10	cool	overcast	normal	true	yes
07-12	mild	sunny	high	false	no
07-14	cool	sunny	normal	false	yes
07-15	mild	rain	normal	false	no
07-20	mild	sunny	normal	true	yes
07-21	mild	overcast	high	true	yes
07-22	hot	overcast	normal	false	yes
07-23	mild	rain	high	true	no
07-26	cool	rain	normal	true	no
07-30	mild	rain	high	false	yes


today      cool      sunny      normal      false      yes

### Nearest Neighbor Classification


Day	Temperature	Outlook	Humidity	Windy	Play Golf?
07-05	hot	sunny	high	false	no
07-06	hot	sunny	high	true	no
07-07	hot	overcast	high	false	yes
07-09	cool	rain	normal	false	yes
07-10	cool	overcast	normal	true	yes
07-12	mild	sunny	high	false	no
07-14	cool	sunny	normal	false	yes
07-15	mild	rain	normal	false	yes
07-20	mild	sunny	normal	true	yes
07-21	mild	overcast	high	true	yes
07-22	hot	overcast	normal	false	yes
07-23	mild	rain	high	true	no
07-26	cool	rain	normal	true	no
12-30	mild	rain	high	false	yes

tomorrow      mild      sunny      normal      false      yes

### Nearest Neighbor Classifier



## Nearest Neighbor Classifier



$k$  nearest neighbors of an example  $x$  are the data points that have the  $k$  smallest distances to  $x$

## Classification

- The predicted class is determined from the nearest neighbor list
- Take the majority vote of class labels among the  $k$ -nearest neighbors

## Distance Functions

- Computes the distance between two examples
  - so that we can find the “nearest neighbor” to a given example
- General Idea:
  - reduce the distance  $d(x_i, x_j)$  of two examples to the distances  $d_A(v_i, v_j)$  between two values for attribute  $A$
- Popular choices
  - Euclidean Distance:**
    - straight-line between two points
$$d(x_1, x_2) = \sqrt{\sum_A d_A(v_{1,A}, v_{2,A})^2}$$
  - Manhattan or City-block Distance:**
    - sum of axis-parallel line segments
$$d(x_1, x_2) = \sum_A d_A(v_{1,A}, v_{2,A})$$


## Distance Functions for Numerical Attributes

- Numerical Attributes:
  - distance between two attribute values
$$d_A(v_i, v_j) = |v_i - v_j|$$
- Normalization:
  - Different attributes are measured on different scales  
→ values need to be normalized in  $[0, 1]$ :
$$\hat{v}_j = \frac{v_j - \min v_j}{\max v_j - \min v_j}$$
- Note:
  - This normalization assumes a (roughly) uniform distribution of attribute values
  - For other distributions, other normalizations might be preferable
    - e.g.: logarithmic for salaries?

## Learning and Decision Trees

- Based on a set of attributes as input, predicted output value, the *decision* is learned
- Boolean or binary classification
  - output values are true or false
  - The simplest case.
- making decisions
  - a sequence of test is performed, testing the value of one of the attributes in each step
  - when a leaf node is reached, its value is returned
  - good correspondence to human decision-making

## Decision tree-induced partition – example



## Learning decision trees


- Goal: Build a **decision tree** to classify examples as positive or negative instances using supervised learning from a training set

- A decision tree** is a tree where

- each non-leaf node has associated with it an attribute (feature)
- each leaf node has associated with it a classification (+ or -)
- each arc has associated with it one of the possible values of the attribute at the node from which the arc is directed

- Generalization: allow for >2 classes

e.g., for stocks, classify into {sell, hold, buy}



## Types of Variables

- Numerical:** Domain is ordered and can be represented on the real line (e.g., age, income)
- Nominal or categorical:** Domain is a finite set without any natural ordering (e.g., occupation, marital status, race)
- Ordinal:** Domain is ordered, but absolute differences between values is unknown (e.g., preference scale, severity of an injury)

## Internal Nodes

- Each internal node has an associated *splitting predicate*. Most common are binary predicates.

Example predicates:

- Age <= 20
- Profession in {student, teacher}
- 5000\*Age + 3\*Salary - 10000 > 0

## Learning decision trees

Problem: decide whether to wait for a table at a restaurant, based on the following attributes:

- Alternate: is there an alternative restaurant nearby?
- Bar: is there a comfortable bar area to wait in?
- Fri/Sat: today Friday or Saturday?
- Hungry: are we hungry?
- Patrons: number of people in the restaurant (None, Some, Full)
- Price: price range (\$, \$\$, \$\$\$)
- Raining: is it raining outside?
- Reservation: have we made a reservation?
- Type: kind of restaurant (French, Italian, Thai, Burger)
- WaitEstimate: estimated waiting time (0-10, 10-30, 30-60, >60)

## Attribute-based representations

- Examples described by attribute values
- E.g., situations where I will/won't wait for a table:

Example	Attributes											Target
	Alt	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Est	Wait	
X <sub>1</sub>	T	F	F	T	Some	\$	F	F	French	0-10	T	
X <sub>2</sub>	T	F	F	T	Full	\$\$	F	T	Thai	30-60	F	
X <sub>3</sub>	F	T	F	F	Some	\$\$	F	F	Burger	0-10	T	
X <sub>4</sub>	T	F	T	T	Full	\$\$	F	T	Thai	10-30	T	
X <sub>5</sub>	T	F	T	F	Full	\$\$\$	F	T	French	>60	F	
X <sub>6</sub>	F	T	F	T	Some	\$\$	T	T	Italian	0-10	T	
X <sub>7</sub>	F	T	F	F	None	\$	T	F	Burger	0-10	F	
X <sub>8</sub>	F	F	F	T	Some	\$\$	T	T	Thai	0-10	T	
X <sub>9</sub>	F	T	T	F	Full	\$\$	T	F	Burger	>60	F	
X <sub>10</sub>	F	T	T	T	Full	\$\$\$	F	T	Italian	10-30	F	
X <sub>11</sub>	F	F	F	F	None	\$	F	F	Thai	0-10	F	
X <sub>12</sub>	T	T	T	T	Full	\$	F	F	Burger	30-60	T	


- Classification of examples is **positive** (T) or **negative** (F)

## Restaurant Sample Set

Example	Attributes										Goal	Explain
	Alt	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Est		
X <sub>1</sub>	Yes	No	No	Yes	Some	\$\$\$	No	Yes	French	0-10	Yes	X <sub>1</sub>
X <sub>2</sub>	Yes	No	No	Yes	Full	\$	No	No	Thai	30-60	No	X <sub>2</sub>
X <sub>3</sub>	No	Yes	No	No	Some	\$	No	No	Burger	0-10	Yes	X <sub>3</sub>
X <sub>4</sub>	Yes	No	Yes	Yes	Full	\$	No	No	Thai	10-30	Yes	X <sub>4</sub>
X <sub>5</sub>	Yes	No	Yes	No	Full	\$\$\$	No	Yes	French	>60	No	X <sub>5</sub>
X <sub>6</sub>	No	Yes	No	Yes	Some	\$	Yes	Yes	Italian	0-10	Yes	X <sub>6</sub>
X <sub>7</sub>	No	Yes	No	No	None	\$	Yes	No	Burger	0-10	No	X <sub>7</sub>
X <sub>8</sub>	No	No	No	Yes	Some	\$\$	Yes	Yes	Thai	0-10	Yes	X <sub>8</sub>
X <sub>9</sub>	No	Yes	Yes	No	Full	\$	Yes	No	Burger	>60	No	X <sub>9</sub>
X <sub>10</sub>	Yes	Yes	Yes	Yes	Full	\$\$\$	Yes	No	Italian	10-30	No	X <sub>10</sub>
X <sub>11</sub>	No	No	No	No	None	\$	No	No	Thai	0-10	No	X <sub>11</sub>
X <sub>12</sub>	Yes	Yes	Yes	Yes	Full	\$	No	No	Burger	30-60	Yes	X <sub>12</sub>

## Decision trees

- One possible representation for hypotheses
- it is a tree for deciding whether to wait:



## Learning Decision Trees

- problem: find a decision tree that agrees with the training set
- trivial solution: construct a tree with one branch for each sample of the training set
  - works perfectly for the samples in the training set
  - may not work well for new samples (generalization)
  - results in relatively large trees
- better solution: find a concise tree that still agrees with all samples
  - corresponds to the simplest hypothesis that is consistent with the training set

## Restaurant Sample Set


Example	Attributes										Goal	Ex
	Alt	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Est		
X1	Yes	No	No	Yes	Some	\$	No	Yes	French	0-10	Yes	X1
X2	Yes	No	No	Yes	Some	\$	No	No	Thai	30-60	No	X2
X3	No	Yes	No	No	Some	\$	No	No	Burger	0-10	Yes	X3
X4	Yes	No	Yes	Yes	Full	\$	No	No	Thai	10-30	Yes	X4
X5	Yes	No	Yes	No	Full	\$\$\$	No	Yes	French	>60	No	X5
X6	No	Yes	No	Yes	Some	\$\$	Yes	Yes	Italian	0-10	Yes	X6
X7	No	Yes	No	No	None	\$	Yes	No	Burger	0-10	No	X7
X8	No	No	No	No	Some	\$\$	Yes	Yes	Thai	0-10	Yes	X8
X9	No	Yes	Yes	No	Full	\$	Yes	No	Burger	>60	No	X9
X10	Yes	Yes	Yes	Yes	Full	\$\$\$	No	Yes	Italian	10-30	No	X10
X11	No	No	No	No	None	\$	No	No	Thai	0-10	No	X11
X12	Yes	Yes	Yes	Yes	Full	\$	No	No	Burger	30-60	Yes	X12

## Restaurant Sample Set

Example	Attributes										Goal	Ex
	Alt	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Est		
X1	Yes	No	No	Yes	Some	\$	No	Yes	French	0-10	Yes	X1
X2	Yes	No	No	Yes	Some	\$	No	No	Thai	30-60	No	X2
X3	No	Yes	No	No	Some	\$	No	No	Burger	0-10	Yes	X3
X4	Yes	No	Yes	Yes	Full	\$	No	No	Thai	10-30	Yes	X4
X5	Yes	No	Yes	No	Full	\$\$\$	No	Yes	French	>60	No	X5
X6	No	Yes	No	Yes	Some	\$\$	Yes	Yes	Italian	0-10	Yes	X6
X7	No	Yes	No	No	None	\$	Yes	No	Burger	0-10	No	X7
X8	No	No	No	Yes	None	\$	Yes	Yes	Thai	0-10	Yes	X8
X9	No	Yes	Yes	No	Full	\$	Yes	No	Burger	>60	No	X9
X10	Yes	Yes	Yes	Yes	Full	\$\$\$	No	Yes	Italian	10-30	No	X10
X11	No	No	No	No	None	\$	No	No	Thai	0-10	No	X11
X12	Yes	Yes	Yes	Yes	Full	\$	No	No	Burger	30-60	Yes	X12


## Choosing an attribute

- Idea: a good attribute splits the examples into subsets that are (ideally) "all positive" or "all negative"



• *Patrons?* is a better choice


## Partial Decision Tree



- Patrons needs further discrimination only for the Full value
- None and Some agree with the Will Wait goal predicate
- the next step will be performed on the remaining samples for the Full value of Patrons

- select best attribute
  - candidate 1: **Pat**      Some and None in agreement with goal
  - candidate 2: **Type**      No values in agreement with goal

## Splitting examples by testing attributes




## Restaurant Sample Set

Example	Attributes	Goal	Exai
X2	Alt Bar Fri Hun Pet Price Rain Res Type Est WillWait		
X4	Yes No No Yes Full \$ No No Thai 30-60 No X2		
X5	Yes No Yes No Full \$\$\$ No Yes French >60 No X5		
X9	No Yes Yes No Full \$ Yes No Burger >60 No X9		
X10	Yes Yes Yes Full \$\$\$ No Yes Italian 10-30 No X10		
X12	Yes Yes Yes Yes Full \$ No No Burger 30-60 Yes X12		


- select next best attribute
  - candidate 1: **Hungry?**      No in agreement with goal
  - candidate 2: **Type?**              No values in agreement with goal

## Partial Decision Tree



- Hungry** needs further discrimination only for the **Yes** value
- No** agrees with the **WillWait** goal predicate
- the next step will be performed on the remaining samples for the **Yes** value of **Hungry**

## Decision Tree Learning



## Finding ‘compact’ decision trees


- Finding the smallest decision tree is difficult.
- There are heuristics that find reasonable decision trees in most practical cases.
- Idea: test the most important attribute first
  - attribute that makes the most difference for the classification of an example
  - hopefully will yield the correct classification with few tests

## Choosing the best attribute

- Key problem: choosing which attribute to split a given set of examples
- Some possibilities are:
  - Random:** Select any attribute at random
  - Least-Values:** Choose the attribute with the smallest number of possible values (e.g. Hungry)
  - Most-Values:** Choose the attribute with the largest number of possible values (e.g. Type)
  - Max-Gain:** Choose the attribute that has the largest expected *information gain*—i.e., attribute that results in smallest expected size of sub-trees rooted at its children
- The ID3 algorithm uses the Max-Gain method of selecting the best attribute

## Choosing an attribute

Idea: a good attribute splits the examples into subsets that are (ideally) “all positive” or “all negative”



Which is better: *Patrons*? or *Type*?


## Decision tree learning

- Aim: find a small tree consistent with the training examples
  - Idea: (recursively) choose "most significant" attribute as root of (sub)tree

```

function DTL(examples, attributes, default) returns a decision tree
  if examples is empty then return default
  else if all examples have the same classification then return the classification
  else if attributes is empty then return MODE(examples)
  else
    best ← CHOOSE-ATTRIBUTE(attributes, examples)
    tree ← a new decision tree with root best
    for each value  $v_i$  of best do
      examples $_i$  ← {elements of examples with  $= v_i$ }
      subtree ← DTL(examples $_i$ , attributes - best, MODE(examples))
      add a branch to tree with label  $v_i$  and subtree subtree
    return tree
  
```

## **Disorder is bad Homogeneity is good**




## **Disorder is bad Homogeneity is good**

- We want a measure that prefers attributes that have a high degree of “order”:
    - Maximum order: All examples are of the same class
    - Minimum order: All classes are equally likely
  - Entropy is a measure for (un-)orderedness
    - all examples of the same class → no information
  - ID3 splits attributes based on their *entropy*.

## Variable Quality Measures

- Let  $S$  be a sample of training instances and  $p_j$  be the proportions of instances of class  $j$  ( $j=1,\dots,J$ ) in  $S$ .
  - Define an **impurity** measure  $I(S)$  that satisfies:
    - $I(S)$  is minimum only when  $p_j=1$  and  $p_j=0$  for  $j \neq i$  (all objects are of the same class);
    - $I(S)$  is maximum only when  $p_j=1/J$  (there is exactly the same number of objects of all classes);
    - $I(S)$  is symmetric with respect to  $p_1,\dots,p_J$ .

## Variable Quality Measures





## Reduction of Impurity: Discrete Variables

- The “best” variable is the variable  $X_i$  that determines a split maximizing the expected reduction of impurity:

$$\Delta I(S, X_i) = I(S) - \sum_j \frac{|S_{xij}|}{|S|} I(S_{xij})$$

where  $S_{xij}$  is the subset of instances from  $S$  s.t.  $X_i = x_{ij}$ .



### Attribution selection criteria:

- Create pure nodes whenever possible
- If pure nodes are not possible, choose the split that leads to the largest decrease in entropy.

$$\text{Entropy} = -P(\text{yes}) \ln[P(\text{yes})] - P(\text{no}) \ln[P(\text{no})]$$

$$\text{General form: Entropy} = -\sum_i P(v_i) \ln[P(v_i)]$$


## Entropy (for two classes)

- $S$  is a set of examples
- $p_{\text{+}}$  is the proportion of examples in class  $\text{+}$
- $p_{\text{-}} = 1 - p_{\text{+}}$  is the proportion of examples in class  $\text{-}$

**Entropy:**

$$E(S) = -p_{\text{+}} \log_2 p_{\text{+}} - p_{\text{-}} \log_2 p_{\text{-}}$$

- Interpretation:**
  - amount of unorderliness in the class distribution of  $S$




## Entropy

- Entropy:**  $H(S) = -p_{(+)} \log_2 p_{(+)} - p_{(-)} \log_2 p_{(-)}$  bits
  - $S$  ... subset of training examples
  - $p_{(+)} / p_{(-}$  ... % of positive / negative examples in  $S$
- Interpretation:** assume item  $X$  belongs to  $S$ 
  - how many bits need to tell if  $X$  positive or negative
- impure (3 yes / 3 no):**

$$H(S) = -\frac{3}{6} \log_2 \frac{3}{6} - \frac{3}{6} \log_2 \frac{3}{6} = 1 \text{ bits}$$
- pure set (4 yes / 0 no):**

$$H(S) = -\frac{4}{4} \log_2 \frac{4}{4} - \frac{0}{4} \log_2 \frac{0}{4} = 0 \text{ bits}$$



© Victor Lavrenko

## Information Gain

- Want many items in pure sets
- Expected drop in entropy after split:

$$\text{Gain}(S, A) = H(S) - \sum_{V \in \text{Values}(A)} \frac{|S_V|}{|S|} H(S_V)$$

$V$  ... possible values of  $A$   
 $S$  ... set of examples  $\{X\}$   
 $S_V$  ... subset where  $X_A = V$

### Mutual Information

- between attribute  $A$  and class labels of  $S$

$$\begin{aligned} \text{Gain}(S, \text{Wind}) &= H(S) - \frac{6}{14} H(S_{\text{weak}}) - \frac{8}{14} H(S_{\text{strong}}) \\ &= 0.94 - \frac{6}{14} * 0.81 - \frac{8}{14} * 1.0 \\ &= 0.049 \end{aligned}$$

$H(S) = 0.94$   
Wind  
Weak Strong  
6 yes / 2 no      3 yes / 3 no  
 $H(S_{\text{weak}}) = 0.81$        $H(S_{\text{strong}}) = 1.0$

© Victor Lavrenko

## Decision tree learning example

10 attributes:

- Alternate:** Is there a suitable alternative restaurant nearby? {yes,no}
- Bar:** Is there a bar to wait in? {yes,no}
- Fri/Sat:** Is it Friday or Saturday? {yes,no}
- Hungry:** Are you hungry? {yes,no}
- Patrons:** How many are seated in the restaurant? {none, some, full}
- Price:** Price level {\$\$,\$\$,\$\$,\$\$}
- Raining:** Is it raining? {yes,no}
- Reservation:** Did you make a reservation? {yes,no}
- Type:** Type of food {French,Italian,Thai,Burger}
- Wait:** {0-10 min, 10-30 min, 30-60 min, >60 min}

## Decision tree learning example


Example	Attributes										Target WillWait
	Alt	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Eat	
$X_1$	T	F	F	T	Some	\$\$\$	F	T	French	0-10	T
$X_2$	T	F	F	T	Full	\$	F	F	Thai	30-60	F
$X_3$	F	T	F	Some	\$	F	F	Burger	0-10	T	
$X_4$	T	F	T	T	Full	\$	F	F	Thai	10-30	T
$X_5$	T	F	T	F	Full	\$\$\$	F	T	French	>60	F
$X_6$	F	T	F	T	Some	\$\$	T	T	Italian	0-10	T
$X_7$	F	T	F	F	None	\$	T	F	Burger	0-10	F
$X_8$	F	F	F	T	Some	\$\$	T	T	Thai	0-10	T
$X_9$	F	T	T	F	Full	\$	T	F	Burger	>60	F
$X_{10}$	T	T	T	T	Full	\$\$\$	F	T	Italian	10-30	F
$X_{11}$	F	F	F	F	None	\$	F	F	Thai	0-10	F
$X_{12}$	T	T	T	T	Full	\$	F	F	Burger	30-60	T

T = True, F = False

$$\text{Entropy} = -\frac{6}{12} \ln \left( \frac{6}{12} \right) - \left( \frac{6}{12} \right) \ln \left( \frac{6}{12} \right) = 0.30$$

6 True,  
6 False

## Decision tree learning example



$$\text{Entropy} = \frac{6}{12} \left[ -\left(\frac{3}{6}\right) \ln \left( \frac{3}{6} \right) - \left(\frac{3}{6}\right) \ln \left( \frac{3}{6} \right) \right] + \frac{6}{12} \left[ -\left(\frac{3}{6}\right) \ln \left( \frac{3}{6} \right) - \left(\frac{3}{6}\right) \ln \left( \frac{3}{6} \right) \right] = 0.30$$

Entropy decrease = 0.30 - 0.30 = 0


## Decision tree learning example



$$\begin{aligned} \text{Entropy} &= \frac{2}{12} \left[ -\left(\frac{0}{2}\right) \ln \left( \frac{0}{2} \right) - \left(\frac{2}{2}\right) \ln \left( \frac{2}{2} \right) \right] + \frac{4}{12} \left[ -\left(\frac{4}{4}\right) \ln \left( \frac{4}{4} \right) - \left(\frac{0}{4}\right) \ln \left( \frac{0}{4} \right) \right] \\ &+ \frac{6}{12} \left[ -\left(\frac{2}{6}\right) \ln \left( \frac{2}{6} \right) - \left(\frac{4}{6}\right) \ln \left( \frac{4}{6} \right) \right] = 0.14 \end{aligned}$$

Entropy decrease = 0.30 - 0.14 = 0.16

## Decision tree learning example



$$\begin{aligned} \text{Entropy} &= \frac{6}{12} \left[ -\left(\frac{3}{6}\right) \ln \left( \frac{3}{6} \right) - \left(\frac{3}{6}\right) \ln \left( \frac{3}{6} \right) \right] + \frac{2}{12} \left[ -\left(\frac{2}{2}\right) \ln \left( \frac{2}{2} \right) - \left(\frac{0}{2}\right) \ln \left( \frac{0}{2} \right) \right] \\ &+ \frac{4}{12} \left[ -\left(\frac{1}{4}\right) \ln \left( \frac{1}{4} \right) - \left(\frac{3}{4}\right) \ln \left( \frac{3}{4} \right) \right] = 0.23 \end{aligned}$$


Entropy decrease = 0.30 - 0.23 = 0.07

## ID3 Algorithm


- A greedy algorithm for decision tree construction developed by Ross Quinlan circa 1987
- Top-down construction of decision tree by recursively selecting “best attribute” to use at the current node in tree
  - Once attribute is selected for current node, generate child nodes, one for each possible value of selected attribute
  - Partition examples using the possible values of this attribute, and assign these subsets of the examples to the appropriate child node
  - Repeat for each child node until all examples associated with a node are either all positive or all negative

## Next step

Given *Patrons* as root node, the next attribute chosen is *Hungry*?



## Decision tree learning example




## ID3-Disadvantages

Error Propagation: Since decision trees work by a series of local decisions, what happens when one of these local decisions is wrong?

- Every decision from that point on may be wrong
- We may never return to the correct path of the tree

## Underfitting and Overfitting (Example)




500 circular and 500 triangular data points.


Circular points:  
 $0.5 \leq \sqrt{x_1^2 + x_2^2} \leq 1$

Triangular points:  
 $\sqrt{x_1^2 + x_2^2} > 0.5$  or  
 $\sqrt{x_1^2 + x_2^2} < 1$

## Decision Trees are Non-linear Classifiers



## Underfitting and Overfitting



**Underfitting:** when model is too simple, both training and test errors are large


## Overfitting

- Definition: If your machine learning algorithm fits noise (i.e. pays attention to parts of the data that are irrelevant) it is overfitting.
- Fact (theoretical and empirical): If your machine learning algorithm is overfitting then it may perform less well on test set data.


## What's Overfitting?

- Overfitting** = Given a hypothesis space  $H$ , a hypothesis  $h \in H$  is said to overfit the training data if there exists some alternative hypothesis  $h' \in H$ , such that
  - $h$  has smaller error than  $h'$  over the training examples, but
  - $h'$  has a smaller error than  $h$  over the entire distribution of instances.


## Example Data Set




## Increasing number of nodes in Decision Trees




Decision Tree with 4 nodes





Decision Tree with 50 nodes




Which tree is better?





## Why Does my Method Overfit ?

- In domains with noise or uncertainty** the system may try to decrease the training error by completely fitting all the training examples



## Model Selection: Using Validation Set

- Divide training data into two parts:
  - Training set:
    - use for model building
  - Validation set:
    - use for estimating generalization error
    - Note: validation set is not the same as test set
- Drawback:
  - Less data available for training

## Avoiding overfitting

- Usually we do not know in advance which are the irrelevant variables
- ...and it may depend on the context
  - For example, if  $y = a \text{ AND } b$  then  $b$  is an irrelevant variable only in the portion of the tree in which  $a=0$

But we can use simple statistics to warn us that we might be overfitting.

## Pruning

- Goal: Prevent overfitting to noise in the data
- Two strategies for “pruning” the decision tree:
  - ◆ *Postpruning* - take a fully-grown decision tree and discard unreliable parts
  - ◆ *Prepruning* - stop growing a branch when information becomes unreliable
- Postpruning preferred in practice—prepruning can “stop too early”

Do not include branches that fit data too specifically

#### Model Selection:

## Incorporating Model Complexity

- Rationale: Occam's Razor
    - Given two models of similar generalization errors, one should prefer the simpler model over the more complex model
    - A complex model has a greater chance of being fitted accidentally by errors in data
    - Therefore, one should include model complexity when evaluating a model
- Gen. Error(Model) = Train. Error(Model, Train. Data) +  $\alpha \times \text{Complexity}(\text{Model})$**

## Estimating Generalization Errors

- Re-substitution errors:** error on training ( $\sum e(t)$ )
- Generalization errors:** error on testing ( $\sum e'(t)$ )
- Methods for estimating generalization errors:
  - Optimistic approach:**  $e'(t) = e(t)$
  - Pessimistic approach:**
    - For each leaf node:  $e'(t) = (e(t)+0.5)$
    - Total errors:  $e'(T) = e(T) + N \times 0.5$  (N: number of leaf nodes)
    - For a tree with 30 leaf nodes and 10 errors on training (out of 1000 instances):
      - Training error =  $10/1000 = 1\%$
      - Generalization error =  $(10 + 30 \times 0.5)/1000 = 2.5\%$
  - Reduced error pruning (REP):**
    - uses validation data set to estimate generalization error

## How to Address Overfitting


### Pre-Pruning (Early Stopping Rule)

- Stop the algorithm before it becomes a fully-grown tree
- Typical stopping conditions for a node:
  - Stop if all instances belong to the same class
  - Stop if all the attribute values are the same
- More restrictive conditions:
  - Stop if number of instances is less than some user-specified threshold
  - Stop if class distribution of instances are independent of the available features (e.g., using  $\chi^2$  test)
  - Stop if expanding the current node does not improve impurity measures (e.g., Gini or information gain).

## Prepruning

- Based on statistical significance test
  - Stop growing the tree when there is no *statistically significant* association between any attribute and the class at a particular node
- Most popular test: *chi-squared test*
- ID3 used chi-squared test in addition to information gain
  - Only statistically significant attributes were allowed to be selected by information gain procedure

## Example of Post-Pruning



## Missing as a separate value

- Missing value denoted “?” in C4.X
- Simple idea: treat missing as a separate value
- Q: When this is not appropriate?**
- A:** When values are missing due to different reasons
  - Example: field **IsPregnant**=missing for a male patient should be treated similarly (no) than for a female patient of age 25 (unknown)

## Handling Missing Attribute Values

- Missing values affect decision tree construction in three different ways:
  - Affects how impurity measures are computed
  - Affects how to distribute instance with missing value to child nodes
  - Affects how a test instance with missing value is classified
- Many possible approaches
  - Treat them as different values
  - Propagate the cases containing such values down the tree without considering them in the "Information Gain" calculation

## Missing values - advanced

- Split instances with missing values into pieces
- A piece going down a branch receives a weight proportional to the popularity of the branch
  - weights sum to 1
  - Info gain works with fractional instances
    - use sums of weights instead of counts
  - During classification, split the instance into pieces in the same way
    - Merge probability distribution using weights

## Computing Impurity Measure

**Before Splitting:**

$$\text{Entropy}(\text{Parent}) = -0.3 \log(0.3) - (0.7) \log(0.7) = 0.8813$$

Tid	Refund	Marital Status	Taxable Income	Class
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	?	Single	90K	Yes

**Split on Refund:**


$$\text{Entropy}(\text{Refund}=\text{Yes}) = 0$$

$$\text{Entropy}(\text{Refund}=\text{No}) = -(2/6)\log(2/6) - (4/6)\log(4/6) = 0.9183$$


$$\text{Entropy}(\text{Children}) = 0.3(0) + 0.6(0.9183) = 0.551$$

$$\text{Gain} = 0.9 \times (0.8813 - 0.551) = 0.3303$$

## Distribute Instances



## Classify Instances



## Numeric attributes

- Standard method: binary splits
  - E.g. temp < 45
- Unlike nominal attributes, every attribute has many possible split points
- Solution is straightforward extension:
  - Evaluate info gain (or other measure) for every possible split point of attribute
  - Choose "best" split point
  - Info gain for best split point is info gain for attribute
- Computationally more demanding

## How to Specify Test Condition?

- Depends on attribute types
  - Nominal
  - Ordinal
  - Continuous
- Depends on number of ways to split
  - 2-way split
  - Multi-way split

## Binary vs. multi-way splits


- Splitting (multi-way) on a nominal attribute exhausts all information in that attribute
  - Nominal attribute is tested (at most) once on any path in the tree
- Not so for binary splits on numeric attributes!
  - Numeric attribute may be tested several times along a path in the tree
- Disadvantage: tree is hard to read
- Remedy:
  - pre-discretize numeric attributes, or
  - use multi-way splits instead of binary ones

## Splitting Based on Nominal Attributes

- **Multi-way split:** Use as many partitions as distinct values.



- **Binary split:** Divides values into two subsets.  
Need to find optimal partitioning.




## Splitting Based on Ordinal Attributes

- **Multi-way split:** Use as many partitions as distinct values.



- **Binary split:** Divides values into two subsets.  
Need to find optimal partitioning.



- What about this split?



## Splitting Based on Continuous Attributes

- Different ways of handling

- **Discretization** to form an ordinal categorical attribute
  - Static – discretize once at the beginning
  - Dynamic – ranges can be found by equal interval bucketing, equal frequency bucketing (percentiles), or clustering.

- **Binary Decision:**  $(A < v)$  or  $(A \geq v)$ 
  - consider all possible splits and finds the best cut
  - can be more compute intensive

## Splitting Based on Continuous Attributes



## Deal with continuous data

- When dealing with nominal data, We evaluated the gain for each possible value
- In continuous data, we have infinite values. What should we do?
  - Continuous-valued attributes may take infinite values, but we have a limited number of values in our instances (at most N if we have N instances)
- Therefore, simulate that you have N nominal values
  - Evaluate information gain for every possible split point of the Attribute Choose the best split point
  - The information gain of the attribute is the information gain of the best split

## Example

### Example

Outlook	Temperature	Humidity	Windy	Play
Sunny	85	85	False	No
Sunny	80	90	True	No
Overcast	83	86	False	Yes
Rainy	75	80	False	Yes
...	...	...	...	...

Continuous attributes

## Split in continuous data

- Split on temperature attribute
 

64	65	68	69	70	71	72	72	75	75	80	81	83	85
Yes	No	Yes	Yes	No	No	Yes	Yes	Yes	No	Yes	Yes	Yes	No
- For example, in the above array of values the split is occurring between 71 and 72 (N distinct values meaning at most N-1 splits)
- Of all such splits , the one with the best Information Gain is chosen for the node

## Splitting Based on GINI

- Used in CART, SLIQ, SPRINT.
- When a node p is split into k partitions (children), the quality of split is computed as,

$$GINI_{split} = \sum_{i=1}^k \frac{n_i}{n} GINI(i)$$

where,  $n_i$  = number of records at child i,  
 $n$  = number of records at node p.

## Measure of Impurity: GINI

- Gini Index for a given node t :

$$GINI(t) = 1 - \sum_j [P(j|t)]^2$$

(NOTE:  $P(j|t)$  is the relative frequency of class j at node t).

- Maximum ( $1 - 1/n_c$ ) when records are equally distributed among all classes, implying least interesting information
- Minimum (0.0) when all records belong to one class, implying most interesting information

<b>C1</b>	<b>0</b>	<b>C1</b>	<b>1</b>	<b>C1</b>	<b>2</b>	<b>C1</b>	<b>3</b>
<b>C2</b>	<b>6</b>	<b>C2</b>	<b>5</b>	<b>C2</b>	<b>4</b>	<b>C2</b>	<b>3</b>
<b>Gini=0.000</b>						<b>Gini=0.444</b>	

## Examples for computing GINI

$$GINI(t) = 1 - \sum_j [P(j|t)]^2$$

<b>C1</b>	<b>0</b>	<b>P(C1) = 0/6 = 0</b>	<b>P(C2) = 6/6 = 1</b>
<b>C2</b>	<b>6</b>	<b>Gini = 1 - (0)^2 - (1)^2 = 1 - 0 - 1 = 0</b>	

$$\begin{array}{|c|c|} \hline C1 & 1 \\ \hline C2 & 5 \\ \hline \end{array} \quad \begin{array}{|c|c|} \hline C1 & 5/6 \\ \hline C2 & 1/6 \\ \hline \end{array}$$


$$P(C1) = 1/6 \quad P(C2) = 5/6$$

$$Gini = 1 - (1/6)^2 - (5/6)^2 = 0.278$$

<b>C1</b>	<b>2</b>	<b>P(C1) = 2/6</b>	<b>P(C2) = 4/6</b>
<b>C2</b>	<b>4</b>	<b>Gini = 1 - (2/6)^2 - (4/6)^2 = 0.444</b>	

## Binary Attributes: Computing GINI Index

- Splits into two partitions
- Effect of Weighing partitions:
  - Larger and Purer Partitions are sought for.



## Categorical Attributes: Computing GINI Index

- For each distinct value, gather counts for each class in the dataset
- Use the count matrix to make decisions

Multi-way split

CarType			
Family	Sports	Luxury	
C1	1	2	1
C2	4	1	1
<b>Gini</b>	<b>0.393</b>		

Two-way split (find best partition of values)

CarType (Sports, Luxury), (Family)			
C1	3	1	
C2	2	4	
<b>Gini</b>	<b>0.400</b>		

CarType (Sports), (Family, Luxury)			
C1	2	2	
C2	1	5	
<b>Gini</b>	<b>0.419</b>		

## Splitting Criteria based on Classification Error

- Classification error at a node t :

$$\text{Error}(t) = 1 - \max_i P(i | t)$$

- Measures misclassification error made by a node.

- Maximum ( $1 - 1/n_i$ ) when records are equally distributed among all classes, implying least interesting information
- Minimum (0.0) when all records belong to one class, implying most interesting information

## Examples for Computing Error

$$\text{Error}(t) = 1 - \max_i P(i | t)$$


C1	<b>0</b>	P(C1) = 0/6 = 0	P(C2) = 6/6 = 1
C2	<b>6</b>	Error = 1 - max (0, 1) = 1 - 1 = 0	

C1	<b>1</b>	P(C1) = 1/6	P(C2) = 5/6
C2	<b>5</b>	Error = 1 - max (1/6, 5/6) = 1 - 5/6 = 1/6	


C1	<b>2</b>	P(C1) = 2/6	P(C2) = 4/6
C2	<b>4</b>	Error = 1 - max (2/6, 4/6) = 1 - 4/6 = 1/3	

## Comparison among Splitting Criteria

For a 2-class problem:



## Misclassification Error vs Gini



$$\text{Gini}(N1) = 1 - (3/3)^2 - (0/3)^2 = 0$$

$$\text{Gini}(N2) = 1 - (4/7)^2 - (3/7)^2 = 0.489$$

$$\text{Gini}(\text{Children}) = \frac{3}{10} * 0 + \frac{7}{10} * 0.489 = 0.342$$

$$\text{Gini}(N2) = 1 - (4/7)^2 - (3/7)^2 = 0.489$$

## Model Evaluation

- Purpose:
  - To estimate performance of classifier on previously unseen data (test set)
- Holdout
  - Reserve k% for training and (100-k)% for testing
  - Random subsampling: repeated holdout
- Cross validation
  - Partition data into k disjoint subsets
  - k-fold: train on k-1 partitions, test on the remaining one
  - Leave-one-out: k=n

## Cross-validation Example

- 3-fold cross-validation

